

Data Science Project

For our analysis and the R programming, we will make use of the following R packages:

```
library(dplyr)    # Easy data cleaning, and the very convenient pipe operator <  
library(ggplot2) # Beautiful plots  
library(mgcv)     # Package to fit generalized additive models
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Loading required package: nlme

Attaching package: 'nlme'

The following object is masked from 'package:dplyr':

collapse

This is mgcv 1.8-40. For overview type 'help("mgcv-package")'.

A basic data science project consists of the following six steps:

1. **State the problem** you are trying to solve. It has to be an unambiguous question that can be answered with data and a statistical or machine learning model. At least, specify: What is being observed? What has to be predicted?
2. **Collect the data**, then clean and prepare it. This is commonly the most time-consuming task, but it has to be done in order to fit a prediction model with the data.
3. **Explore the data**. Get to know its properties and quirks. Check numerical summaries of your metric variables, tables of the categorical data, and plot univariate and multivariate representations of your variables. By this, you also get an overview of the quality of the data and can find outliers.
4. **Check if any variables may need to be transformed**. Most commonly, this is a logarithmic transformation of skewed measurements such as concentrations or times. Also, some variables might have to be split up into two or more variables.

5. **Choose a model and train it on the data.** If you have more than one candidate model, apply each and evaluate their goodness-of-fit using independent data that was not used for training the model.
6. **Use the best model to make your final predictions.**

```
data(rivers)
ls()

'avgpm25' · 'awake_only' · 'filename' · 'msleep' · 'msleep10' · 'msleep5' · 'mtcars' · 'mtcars_data' ·
'nlants_data' · 'rivers' · 'sleenData' · 'url'

data()

summary(rivers)
sd(rivers)

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
135.0   310.0   425.0   591.2   680.0  3710.0
493.87084203459

dim(rivers)

NULL

getwd()
dir()

'/content'
'avgpm25.csv' · 'sample_data'

avgpm25<-read.csv('avgpm25.csv')
class(avgpm25)
dim(avgpm25)

'data.frame'
576 · 5

str(avgpm25)

'data.frame':  576 obs. of  5 variables:
 $ pm25      : num  9.77 9.99 10.69 11.34 12.12 ...
 $ fips      : int  1003 1027 1033 1049 1055 1069 1073 1089 1097 1103 ...
 $ region    : chr   "east" "east" "east" "east" ...
 $ longitude : num  -87.7 -85.8 -87.7 -85.8 -86 ...
 $ latitude  : num   30.6 33.3 34.7 34.5 34 ...
< table of extent 0 >

nrow(avgpm25)
ncol(avgpm25)
```

```
object.size(avgpm25)
names(avgpm25)
head(avgpm25)
summary(avgpm25)
table(avgpm25$region)

576
5
22088 bytes
'pm25' · 'fips' · 'region' · 'longitude' · 'latitude'
A data.frame: 6 × 5

  pm25    fips region longitude latitude
  <dbl> <int> <chr>    <dbl>    <dbl>
1  9.771185  1003  east   -87.74826  30.59278
2  9.993817  1027  east   -85.84286  33.26581
3 10.688618  1033  east   -87.72596  34.73148
4 11.337424  1049  east   -85.79892  34.45913
5 12.119764  1055  east   -86.03212  34.01860
6 10.827805  1069  east   -85.35039  31.18973

  pm25      fips      region      longitude
Min.   : 3.383  Min.   : 1003  Length:576  Min.   : -158.04
1st Qu.: 8.549  1st Qu.:16038  Class :character 1st Qu.: -97.38
Median :10.047  Median :28034  Mode  :character Median : -87.37
Mean   : 9.836  Mean   :28431           Mean   : -91.65
3rd Qu.:11.356 3rd Qu.:41045           3rd Qu.: -80.72
Max.   :18.441  Max.   :56039           Max.   : -68.26

  latitude
Min.   :19.68
1st Qu.:35.30
Median :39.09
Mean   :38.56
3rd Qu.:41.75
Max.   :64.82

east west
442  134

install.packages('downloader')

Installing package into ‘/usr/local/lib/R/site-library’
(as ‘lib’ is unspecified)
```

```
library(downloader)
```

```
url <- "https://raw.githubusercontent.com/genomicsclass/dagdata/master/inst/extdata/msleep_
filename <- "msleep_ggplot2.csv"
if (!file.exists(filename)) download(url,filename)
msleep <- read.csv("msleep_ggplot2.csv")
```

```
names(msleep)

'name' · 'genus' · 'vore' · 'order' · 'conservation' · 'sleep_total' · 'sleep_rem' · 'sleep_cycle' · 'awake' ·
'brainwt' · 'bodywt'

head(msleep,2)
tail(msleep,2)
```

A data.frame: 2 × 11

	name	genus	vore	order	conservation	sleep_total	sleep_rem	sleep_cyc
	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>
1	Cheetah	Acinonyx	carni	Carnivora	lc	12.1	NA	1
2	Owl monkey	Aotus	omni	Primates	NA	17.0	1.8	1

A data.frame: 2 × 11

	name	genus	vore	order	conservation	sleep_total	sleep_rem	sleep_cycle
	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>
82	Arctic fox	Vulpes	carni	Carnivora	NA	12.5	NA	NA

```
str(msleep)

'data.frame': 83 obs. of 11 variables:
 $ name      : chr  "Cheetah" "Owl monkey" "Mountain beaver" "Greater short-tailed
 $ genus     : chr  "Acinonyx" "Aotus" "Aplodontia" "Blarina" ...
 $ vore      : chr  "carni" "omni" "herbi" "omni" ...
 $ order     : chr  "Carnivora" "Primates" "Rodentia" "Soricomorpha" ...
 $ conservation: chr  "lc" NA "nt" "lc" ...
 $ sleep_total : num  12.1 17 14.4 14.9 4 14.4 8.7 7 10.1 3 ...
 $ sleep_rem  : num  NA 1.8 2.4 2.3 0.7 2.2 1.4 NA 2.9 NA ...
 $ sleep_cycle : num  NA NA NA 0.133 0.667 ...
 $ awake     : num  11.9 7 9.6 9.1 20 9.6 15.3 17 13.9 21 ...
 $ brainwt   : num  NA 0.0155 NA 0.00029 0.423 NA NA NA 0.07 0.0982 ...
 $ bodywt    : num  50 0.48 1.35 0.019 600 ...
```

```
awake_only<-msleep$awake
head(awake_only,20)
summary(msleep$awake)
table(msleep$conservation)
```

11.9 · 7 · 9.6 · 9.1 · 20 · 9.6 · 15.3 · 17 · 13.9 · 21 · 18.7 · 14.6 · 14 · 11.5 · 13.7 · 15.7 · 14.9 · 6.6 · 18.7 · 6
Min. 1st Qu. Median Mean 3rd Qu. Max.
4.10 10.25 13.90 13.57 16.15 22.10

cd domesticated en lc nt vu
2 10 4 27 4 7

```
summary(msleep)
```

name	genus	vore	order
Length:83	Length:83	Length:83	Length:83
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character

conservation	sleep_total	sleep_rem	sleep_cycle
Length:83	Min. : 1.90	Min. :0.100	Min. :0.1167
Class :character	1st Qu.: 7.85	1st Qu.:0.900	1st Qu.:0.1833
Mode :character	Median :10.10	Median :1.500	Median :0.3333
	Mean :10.43	Mean :1.875	Mean :0.4396
	3rd Qu.:13.75	3rd Qu.:2.400	3rd Qu.:0.5792
	Max. :19.90	Max. :6.600	Max. :1.5000
		NA's :22	NA's :51

awake	brainwt	bodywt
Min. : 4.10	Min. :0.00014	Min. : 0.005
1st Qu.:10.25	1st Qu.:0.00290	1st Qu.: 0.174
Median :13.90	Median :0.01240	Median : 1.670
Mean :13.57	Mean :0.28158	Mean : 166.136
3rd Qu.:16.15	3rd Qu.:0.12550	3rd Qu.: 41.750
Max. :22.10	Max. :5.71200	Max. :6654.000
	NA's :27	

```
msleep10 <- select(msleep, -sleep_cycle)
```

```
dim(msleep10)
names(msleep10)
```

```
83 · 10
'name' · 'genus' · 'vore' · 'order' · 'conservation' · 'sleep_total' · 'sleep_rem' · 'awake' · 'brainwt' · 'bodywt'
```

```
msleep5 <- select(msleep, -c('name','genus','vore','order','conservation','sleep_cycle'))
```

```
dim(msleep5)
names(msleep5)
```

```
83 · 5
'sleep_total' · 'sleep_rem' · 'awake' · 'brainwt' · 'bodywt'
```

```
sleepData <- select(msleep, name, sleep_total)
```

```
str(msleep)
```

```
'data.frame': 83 obs. of 11 variables:
 $ name      : chr  "Cheetah" "Owl monkey" "Mountain beaver" "Greater short-tailed
 $ genus     : chr  "Acinonyx" "Aotus" "Aplodontia" "Blarina" ...
 $ vore      : chr  "carni" "omni" "herbi" "omni" ...
 $ order     : chr  "Carnivora" "Primates" "Rodentia" "Soricomorpha" ...
 $ conservation: chr  "lc" NA "nt" "lc" ...
 $ sleep_total : num  12.1 17 14.4 14.9 4 14.4 8.7 7 10.1 3 ...
 $ sleep_rem  : num  NA 1.8 2.4 2.3 0.7 2.2 1.4 NA 2.9 NA ...
 $ sleep_cycle : num  NA NA NA 0.133 0.667 ...
 $ awake     : num  11.9 7 9.6 9.1 20 9.6 15.3 17 13.9 21 ...
```

5/7/22, 1:57 AM28_Experiment5.ipynb - Colaboratory

```
$ brainwt      : num  NA 0.0155 NA 0.00029 0.423 NA NA NA 0.07 0.0982 ...
$ bodywt       : num   50 0.48 1.35 0.019 600 ...
'data.frame':   83 obs. of  2 variables:
 $ name        : chr   "Cheetah" "Owl monkey" "Mountain beaver" "Greater short-tailed s
 $ sleep_total: num   12.1 17 14.4 14.9 4 14.4 8.7 7 10.1 3 ...
```

```
str(sleepData)

'data.frame':   83 obs. of  2 variables:
 $ name        : chr   "Cheetah" "Owl monkey" "Mountain beaver" "Greater short-tailed s
 $ sleep_total: num   12.1 17 14.4 14.9 4 14.4 8.7 7 10.1 3 ...
```

```
head(select(msleep, -c('name','bodywt','vore')))
```


A data.frame: 6 × 8

	genus	order	conservation	sleep_total	sleep_rem	sleep_cycle	awake
	<chr>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	Acinonyx	Carnivora	lc	12.1	NA	NA	11.9
2	Aotus	Primates	NA	17.0	1.8	NA	7.0
3	Aplodontia	Rodentia	nt	14.4	2.4	NA	9.6
4	Blarina	Soricomorpha	lc	14.9	2.3	0.1333333	9.1
5	Bos	Artiodactyla	domesticated	4.0	0.7	0.6666667	20.0
6	Bradypus	Pilosa	NA	14.4	2.2	0.7666667	9.6

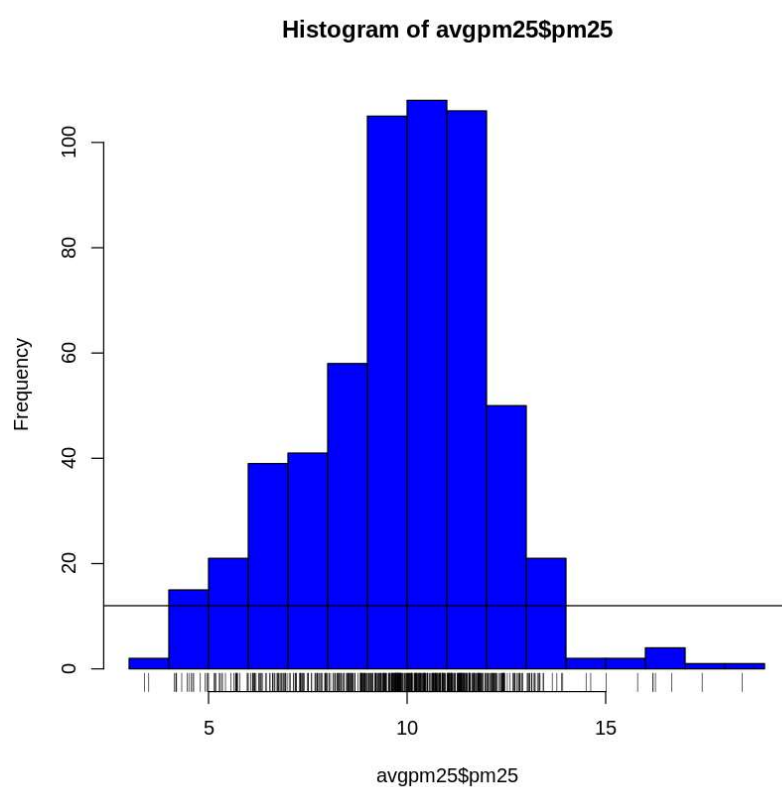
```
ls()
summary(avgpm25$pm25)

'avgpm25' · 'awake_only' · 'filename' · 'msleep' · 'msleep10' · 'msleep5' · 'mtcars' · 'mtcars_data' ·
'plants_data' · 'rivers' · 'sleepData' · 'url'
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
  2.282    8.549   10.047    9.826   11.256   18.441


boxplot(avgpm25$pm25, col='red')
abline(h=12)
```



```
hist(avgpm25$pm25,col='blue')  
rug(avgpm25$pm25)
```




```
hist(avgpm25$pm25,col='blue')  
rug(avgpm25$pm25)  
abline(v=12, col='red')
```




```
barplot(table(avgpm25$region))  
table(avgpm25$region)
```

```
east west  
442  134
```



```
barplot(table(msleep$conservation), main="Title")  
table(msleep$conservation)
```

✓ 2s completed at 01:57

×